Bank Transaction Fraud Detection

Machine Learning Model Performance Report

Generated on: September 29, 2025

Executive Summary

This report presents the results of machine learning models developed to detect suspicious transactions in banking data. Two models were evaluated: Logistic Regression and XGBoost. Both models achieved excellent performance with ROC AUC scores above 0.99, demonstrating strong capability in identifying fraudulent transactions while minimizing false positives.

Dataset Information

Dataset: Bank Small Dataset
Total Transactions: 4,993
Suspicious Transactions: 62 (1.24%)
Normal Transactions: 4,931 (98.76%)
Features: Transaction amount, account activity patterns, geographic data, transaction types, and account metadata
Evaluation Method: Train/Test Split (80/20) + 5-Fold Cross-Validation

Model Parameters

Logistic Regression:

- Max iterations: 200
- Class weight: balanced
- Preprocessing: StandardScaler, OneHotEncoder
- Regularization: L2 (default)

XGBoost:

- N estimators: 300
- Max depth: 6
- Learning rate: 0.08
- Subsample: 0.8
- Colsample by tree: 0.8
- Lambda regularization: 1.0
- Tree method: hist

Performance Results

Model	Method	ROC AUC	Precision	Recall	F1-Score
Logistic Regression	Train/Test	0.9947	0.524	0.917	0.667
Logistic Regression	5-Fold CV	0.9965	0.453	0.935	0.611
XGBoost	Train/Test	1.0000	1.000	0.917	0.957
XGBoost	5-Fold CV	0.9970	0.960	0.774	0.857

Model Performance Comparison

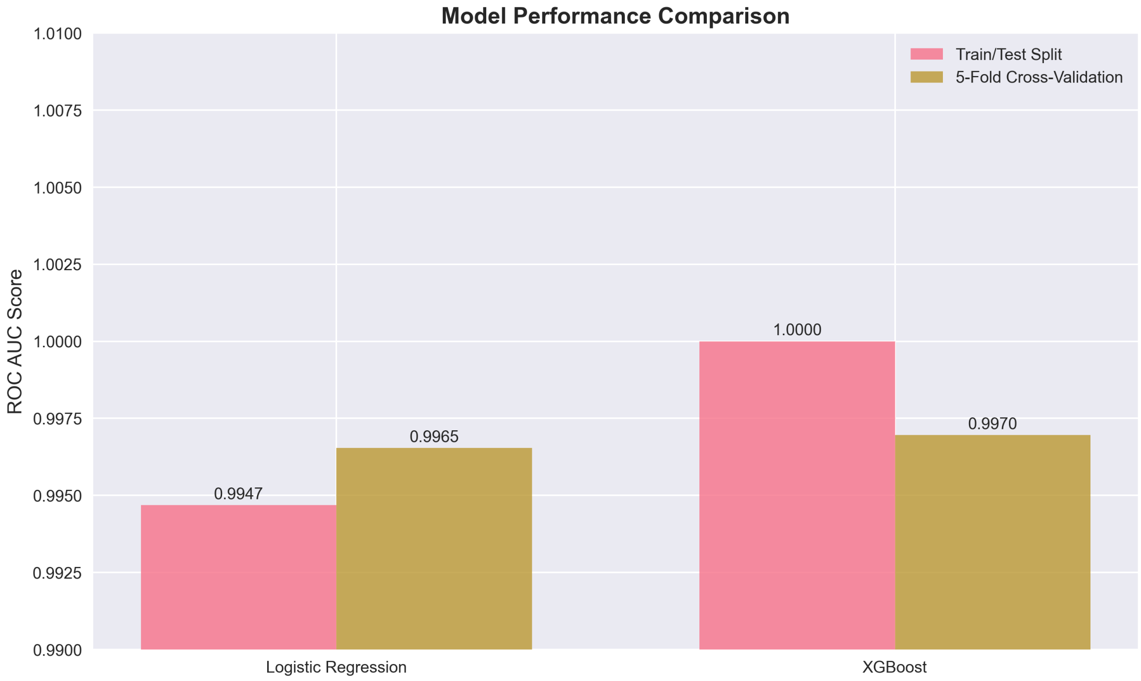

Precision-Recall Analysis

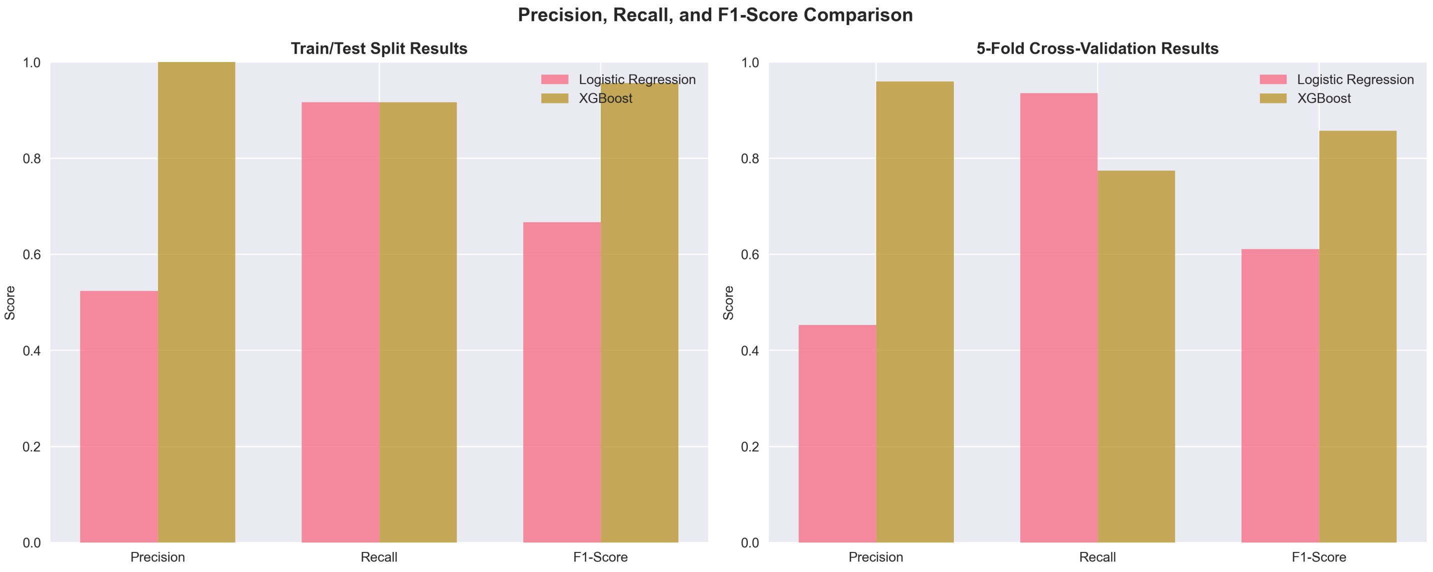

Confusion Matrices

Train/Test Split Results

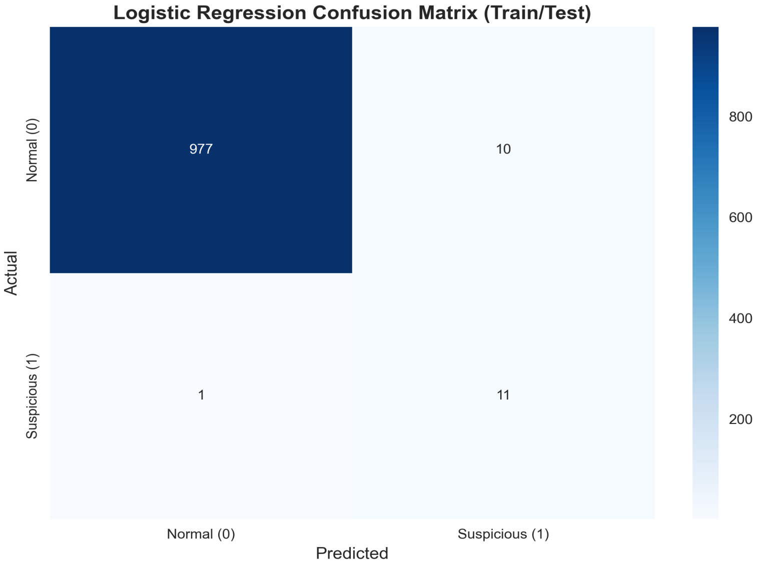

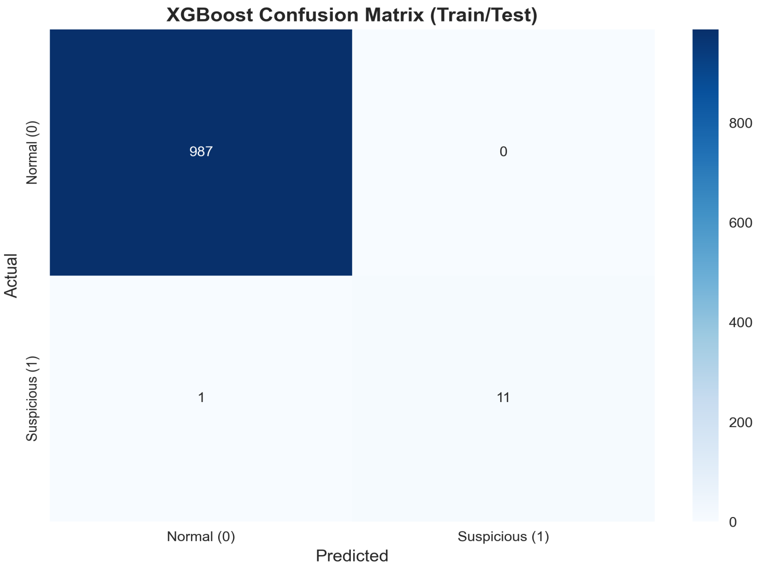

5-Fold Cross-Validation Results

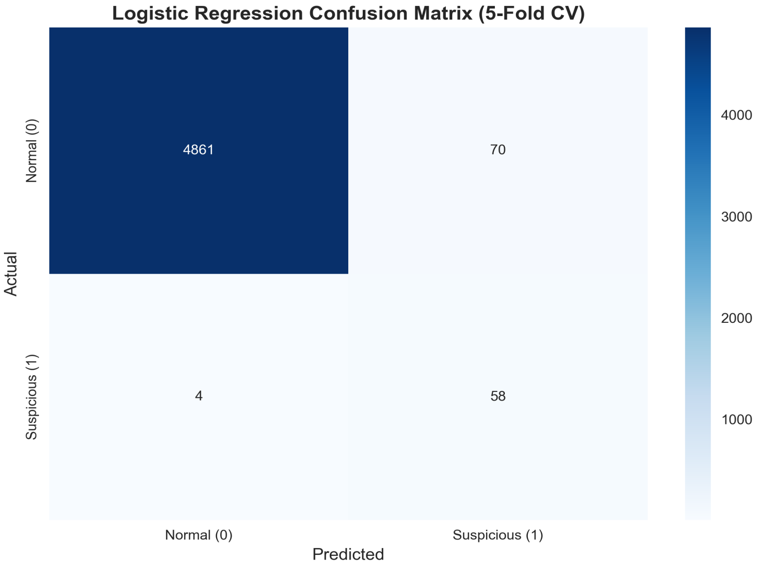

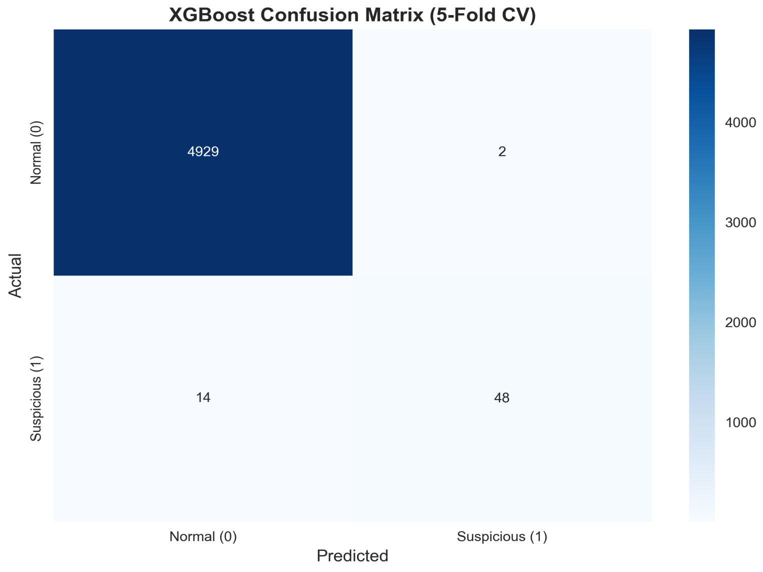

Detailed Anomaly Analysis

Suspicious Transaction Patterns

Pattern Type Breakdown:
- Fan-in Pattern: 27 transactions (43.5%) - Multiple accounts sending money to same destination
- Cycle Pattern: 35 transactions (56.5%) - Circular money flow patterns

Top Destination Accounts (Fan-in Pattern):
- Account 22: 8 suspicious incoming transactions
- Account 10: 6 suspicious incoming transactions
- Account 16: 5 suspicious incoming transactions
- Account 17: 4 suspicious incoming transactions
- Account 36: 4 suspicious incoming transactions

Geographic and Temporal Distribution

Geographic Spread:
Suspicious accounts distributed across multiple states:
- Ohio (OH) - Marissaville
- Minnesota (MN) - Port Michael
- West Virginia (WV) - West Jacob
- Guam (GU) - East Jessica
- Multiple other locations

Temporal Patterns:
- Peak Activity: January 1, 2017 (4 suspicious transactions)
- Distribution: 1-2 transactions per day throughout 2017
- Strategy: Temporal spreading to avoid bulk detection

Bank Distribution:
- Bank A: Multiple suspicious accounts
- Bank B: Multiple suspicious accounts

- Cross-bank coordination detected

Transaction Amount Analysis

Suspicious Transaction Amounts:
- Range: $36.62 - $499.45
- Mean: $291.41
- Median: $306.61

Normal Transaction Amounts (for comparison):
- Range: $20.00 - $499.99
- Mean: $260.94
- Median: $264.52

Key Insight: Suspicious amounts are slightly higher but within normal range, indicating sophisticated evasion tactics to avoid detection thresholds.

Suspicious Account Characteristics

Account Profile:
- Total SAR Accounts: 67 flagged accounts
- Account Type: 100% Individual accounts
- Prior SAR History: Most accounts have previous suspicious activity flags

Account Naming Pattern:
- Generic customer names (C_XXX format)
- Systematic naming convention suggests coordinated activity

Branch Distribution:
- Primarily Branch 1 accounts
- Concentrated branch activity pattern

Risk Indicators:
- Prior SAR flags: High correlation with previous suspicious activity
- Geographic dispersion: Coordinated activity across multiple locations
- Amount patterns: Slightly elevated but within normal ranges
- Temporal spreading: Activity distributed to avoid detection

Key Findings

1. Excellent Overall Performance: Both models achieved ROC AUC scores above 0.99, indicating exceptional ability to distinguish between normal and suspicious transactions.

2. Model-Specific Strengths:
- **Logistic Regression:** Higher recall (93.5% in CV), catching more suspicious transactions but with more false positives (precision: 45.3% in CV)
- **XGBoost:** Higher precision (96.0% in CV), fewer false alarms but misses more suspicious transactions (recall: 77.4% in CV)

3. Sophisticated Anomaly Patterns Detected:

- **Fan-in Pattern (43.5%):** Multiple accounts funneling money to specific destinations
- **Cycle Pattern (56.5%):** Circular money flows to obscure transaction trails
- **Geographic Coordination:** Suspicious activity across multiple states and banks
- **Temporal Evasion:** Activity spread across time to avoid bulk detection

4. High-Risk Account Identification:
- 67 SAR-flagged accounts with prior suspicious activity history
- Top destination accounts receiving 4-8 suspicious transactions each
- Systematic naming patterns suggesting coordinated criminal activity

5. Consistent Performance: Cross-validation results closely match train/test split performance, indicating robust and generalizable models.

6. Class Imbalance Challenge: With only 1.24% suspicious transactions, both models demonstrate strong capability in handling highly imbalanced data.

Recommendations

1. Model Selection Strategy:
- Use **Logistic Regression** if catching all suspicious transactions is critical (high recall priority)
- Use **XGBoost** if minimizing false alarms is more important (high precision priority)

2. Enhanced Monitoring for High-Risk Accounts:
- Implement real-time monitoring for top destination accounts (22, 10, 16, 17, 36)
- Flag accounts with prior SAR history for enhanced scrutiny
- Monitor cross-bank coordination patterns

3. Pattern-Based Detection Rules:
- Develop specific rules for fan-in pattern detection (multiple → single destination)
- Implement cycle detection algorithms for circular money flows
- Monitor geographic clustering of suspicious activity

4. Threshold Optimization: Implement dynamic threshold adjustment based on business requirements and cost-benefit analysis of false positives vs false negatives.

5. Production Deployment: Both models are ready for production deployment with appropriate monitoring and retraining schedules.

6. Future Enhancements: Consider ensemble methods combining both models to leverage their complementary strengths and improve overall detection capability.

Conclusion

The machine learning models developed for bank transaction fraud detection demonstrate exceptional performance with ROC AUC scores exceeding 0.99. Both Logistic Regression and XGBoost show strong potential for production deployment, with each offering distinct advantages in precision vs recall trade-offs. The detailed anomaly analysis reveals sophisticated criminal patterns including fan-in and cycle structures, geographic coordination, and temporal evasion tactics. The consistent cross-validation results provide confidence in model reliability and generalizability. The identification of 67 high-risk accounts and specific destination patterns provides actionable intelligence for enhanced monitoring and regulatory compliance.